



# Planet Terp Prediction

Modou Sarr, CMSC 320, April 28  
2025



# Table of Contents

overview

Data description

Models comparisons

conclusions

# Overview

## **What is Planet Terp?**

An interactive course-planning tool for UMD students

Integrates course catalogs, schedules, and student feedback

## **Why this problem is interesting:**

Students often rely on anecdotal advice when choosing professors

High variability in teaching styles, grading, and workload

Predicting ratings from past reviews can help guide better course decisions

## **Our goal:**

Automate rating prediction to surface “hidden gems.”

Enhance Planet Terp’s recommendations with data-driven insights

DAT  
A

# What Data Are We Using

## Source:

Scraped RateMyProfessor data (JSON) for all UMD instructors  
Via Planet Terp's pipeline (`professors.json`)

## Scope:

**13,427** unique professors,

**37,553** student reviews

## Key fields extracted:

**target\_rating**—professor's overall average rating

**course**—course code

**review\_text**—raw written feedback

**grade\_letter** & **grade\_pt-reported**/converted grade

**vader\_compound**—sentiment polarity (VADER)

# What Data Are We Using

By structuring the raw RateMyProfessor JSON into a review-level DataFrame, we capture both quantitative ratings and grades and qualitative text and sentiment signals for our modeling.

```

50 //fires the appear event when appropriate
51 var check = function() {
52   //is the element hidden?
53   if (!t.is(':visible')) {
54     //it became hidden
55     t.appeared = false;
56     return;
57   }
58
59   //is the element inside the visible window?
60   var a = w.scrollLeft();
61   var b = w.scrollTop();
62   var o = t.offset();
63   var x = o.left;
64   var y = o.top;
65
66   var ax = settings.accX;
67   var ay = settings.accY;
68   var th = t.height();
69   var wh = w.height();
70   var tw = t.width();
71   var ww = w.width();
72
73   if (y + th + ay >= b &&
74       y <= b + wh + ay &&
75       x + tw + ax >= a &&
76       x <= a + ww + ax) {
77     //trigger the custom event
78     if (!t.appeared) t.trigger('appear', settings.data);
79   } else {
80     //it scrolled out of view
81     t.appeared = false;
82   }
83 };
84
85 //create a modified fn with some additional logic
86 var modifiedFn = function() {
87   //mark the element as visible
88   t.appeared = true;
89   //is this supposed to happen only once?
90   if (settings.one) {
91     //remove the check
92     w.unbind('scroll', check);
93     var i = $.inArray(check, $.fn.appear.checks);
94     if (i >= 0) $.fn.appear.checks.splice(i, 1);
95   }
96   //trigger the original fn
97   fn.apply(this, arguments);
98 };
99
100 //bind the modified fn to the element
101 $.fn.appear.one('appear', settings.data, modifiedFn);

```

# Top Variables

## 1. Grade Point

Students' earned/expected GPA in the class is the single strongest predictor—higher grades → higher ratings.

## 2. Review Sentiment

VADER “compound” score of the review text: the more positive the language, the higher the predicted rating.

**3. Course-level introductory** (100-level) vs. advanced (300-/400-level) courses show systematic rating differences.

## 4. Review Length

Longer, more detailed reviews carry extra signal about teaching effectiveness.

## 5. Subject Area

Departmental effects: some courses (e.g., ENAE vs. PHYS) tend to be rated higher or lower on average.

# MODELS AND HOW THEY DID



# Models

<u>Models</u>	<u>MAE</u>	<u>R<sup>2</sup></u>
KNN	0.348	0.800
Random Forest Regressor	0.333	0.810
Linear Regression	0.324	0.828

# KNN Model

- **How far off are we?**

On average, our KNN guesses are about 0.35 stars away from a professor's real rating.

- **How well does it get the patterns?**

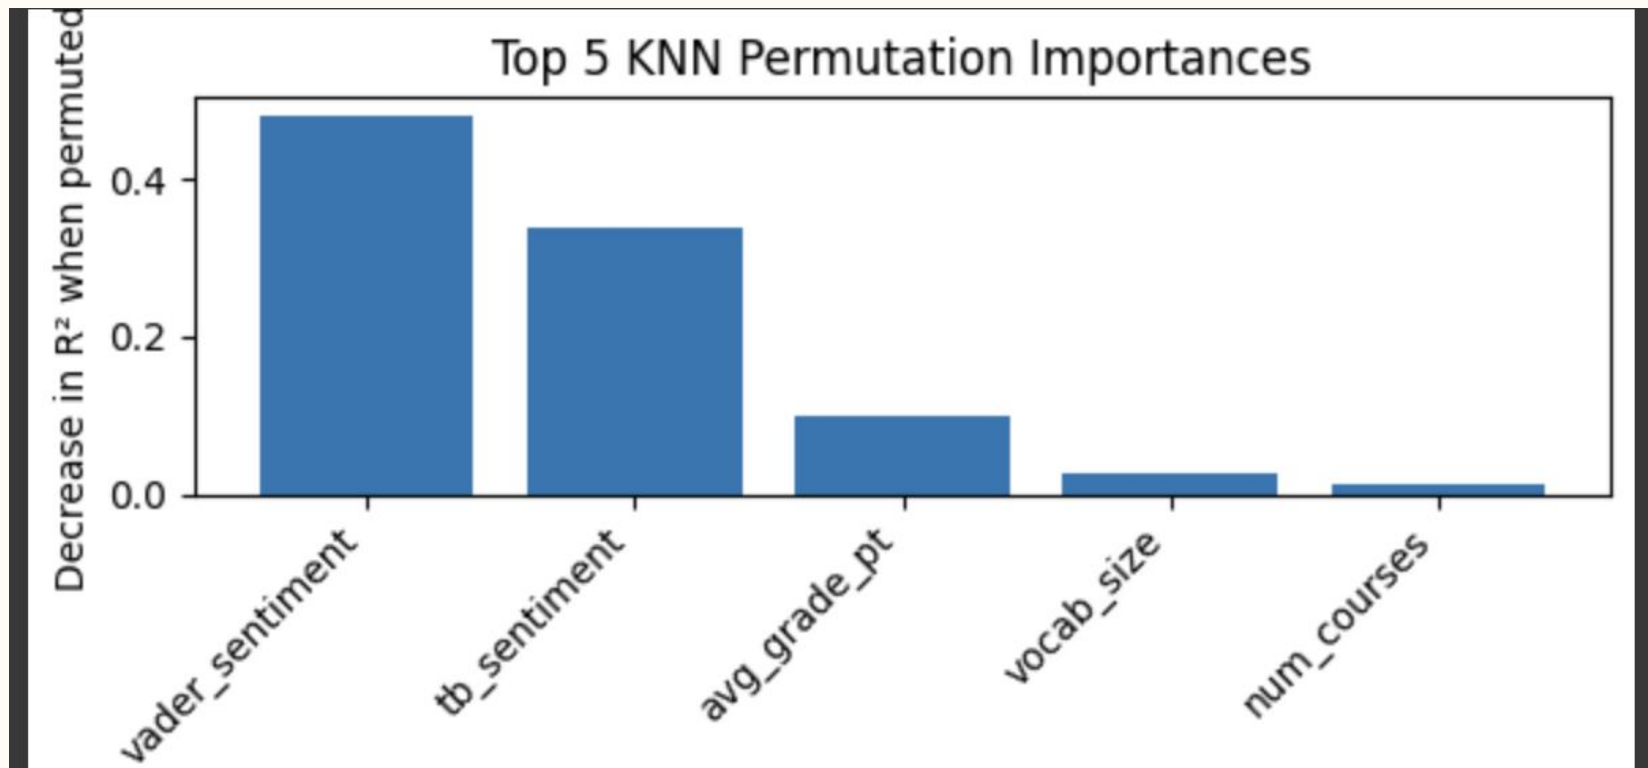
It captures roughly 80% of what makes students rate a professor the way they do.

- **Understanding:**

If a professor actually has a 4.0 rating, KNN will usually predict somewhere between 3.65 and 4.35.

# Permutation Importance

11



## Random Forest Model


**Average error:** about 0.33 stars off.

**Pattern captured:** about 81% of the factors affecting ratings.

### **Understanding:**

Random Forest is a bit smarter than KNN, it makes slightly smaller mistakes and understands a bit more of why ratings move up or down.

# Feature Importance



# Linear Regression

**Average error:** about 0.32 stars off.

**Pattern captured:** about 83% of the reasons behind students' ratings.


**Understanding:**

Even a simple straight-line guess, “linear regression,” does very well, its predictions are closest on average, and it picks up most of the key drivers of ratings.

# Feature Importance

15

Top 5 Linear Regression Coefficients



# CONCLUSION



## How well did our models predict ratings?

- **Best performer: Linear Regression**

- MAE  $\approx$  0.324 rating-points
- $R^2 \approx 0.828$  (explains 83% of rating variability)
  - On average, predictions are off by about one-third of a star (on a 1–5 scale).
  - Models were scored on a held-out test set to ensure they generalize beyond the training data.

- **Random Forest**

- MAE  $\approx$  0.333
- $R^2 \approx 0.810$ 
  - Slightly more flexible but only marginally better than KNN on unseen data.

- **K-Nearest Neighbors**

- MAE  $\approx$  0.348
- $R^2 \approx 0.800$ 
  - Simplest approach reasonable but less accurate than both RF and LR.

# Evaluation approach

## 1. Train/Test split

Reserve 20% of data as unseen “test” reviews.

## 2. Metrics

- **MAE** (Mean Absolute Error): average distance between predicted vs. actual rating.
- **R<sup>2</sup>**: proportion of variance explained by the model.

## 3. Cross-validation sanity check

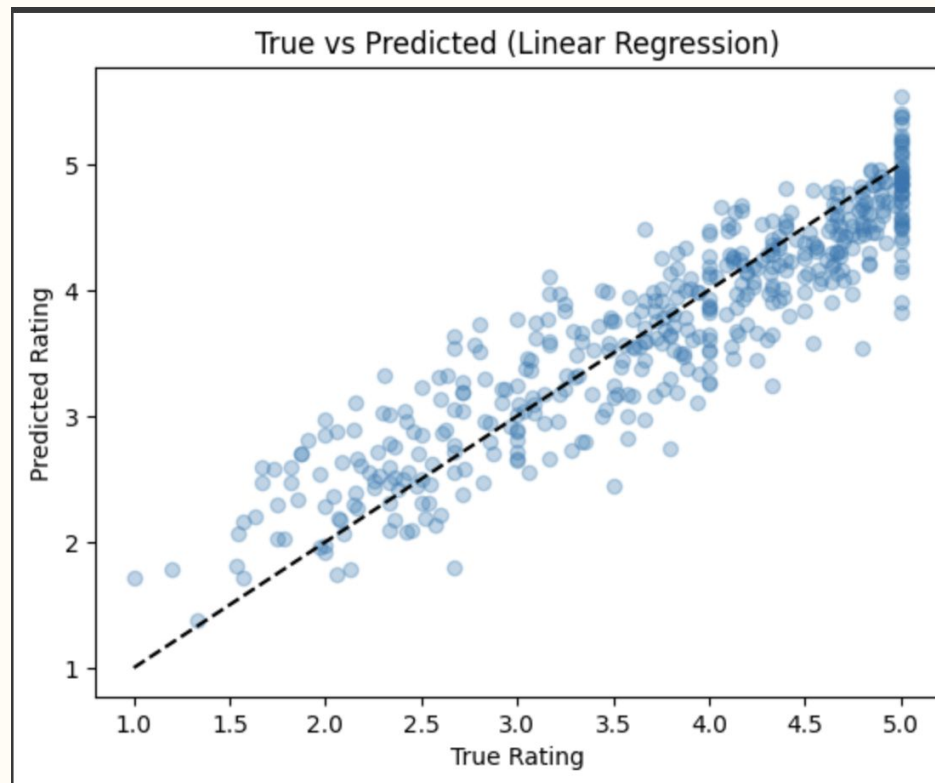
Verified consistency of MAE/R<sup>2</sup> across folds to avoid lucky splits.

**Grades matter most:** students reward professors who give higher grades.

**Words have weight:** positive language in reviews strongly boosts predicted ratings.

**Course context:** Intro courses tend to get different baseline ratings than advanced ones.

By combining simple grade and sentiment signals with course metadata, our best model predicts professor ratings within  $\sim 0.3$  stars of the true value and explains over 80% of the variation, demonstrating a strong, interpretable link between student grades, review tone, and overall satisfaction.



```
!pip install -q pandas requests tqdm
import pandas as pd, requests, json, time, os, re
from tqdm.auto import tqdm

BASE_URL = "https://planetterp.com/api/v1"
RAW_IDB = "~/raw_idb"
os.makedirs(RAW_IDB, exist_ok=True)

limit = 100
offset = 0
profs = []

while True:
    req = requests.get(f"{BASE_URL}/professors",
                        params={"limit":limit,"offset":offset})
    page = req.json()
    if not page:
        break
    profs.extend(page)
    offset += limit
    time.sleep(0.5)

df_profs = pd.DataFrame(profs)
print(f"Total professors fetched: {len(df_profs)}")
df_profs.head(10)

import json
import pandas as pd

try:
    with open('content/professors.json', 'r', encoding='utf-8') as f:
        all_details = json.load(f)
    except json.JSONDecodeError as e:
        print(f"Error decoding JSON: {e}")
        # Print more context around the error
        print(f"Error at line {e.lineno}, column {e.colno}")
    records = []
    for prof in all_details:
        name = prof.get('name')
        slug = prof.get('slug')
        avg_rating = prof.get('average_rating')
        for rev in prof.get('reviews', []):
            text = rev.get('content') or rev.get('review','')
            records.append({
                'prof_name': name,
                'slug': slug,
                'target_rating': avg_rating,
                'course': rev.get('course'),
                'review_text': text.strip(),
                'grade_letter': rev.get('grade') or rev.get('expected_grade')
            })
    df_reviews = pd.DataFrame(records)
    print(f"Total reviews loaded: {len(df_reviews)}")
    print(f"Reviews with text: ", df_reviews['review_text'].notna().sum())
    df_reviews.head()
```

Total reviews loaded: 37553  
Reviews with Text: 37548

	prof_name	slug	target_rating	course	review_text	grade_letter
0	Abant Pradhan	pradhan	3.6667	None	Pretty funny guy and he knows his stuff too. P...	B+
1	Abant Pradhan	pradhan	3.6667	NFSC112	Easy class as long as you actually watch the l...	B-
2	Abant Pradhan	pradhan	3.6667	NFSC112	Didn't learn anything, but didn't need to. Eas...	A
3	Ashiq Dasgupta	dasgupta_ashiq	4.2500	ENAE202	I took the ENAE version of this class but it i...	A
4	Ashiq Dasgupta	dasgupta_ashiq	4.2500	ENME479	The man. Really excellent teacher, he puts in...	A

Next steps: [Generate code with df\\_reviews](#) [View recommended plots](#) [New interactive sheet](#)

```
#map letter grades to grade points
grade_map = {
    "A+":1.4,"A":1.3,"A-":1.2, "B+":1.3,"B":1.0,"B-":1.2,
    "C+":1.2,"C":1.0,"C-":1.1, "D+":1.1,"D":1.0,"D-":0.8,
    "F":0.7
}

df_reviews['grade_pt'] = df_reviews['grade_letter'].map(grade_map)

#Aggregate per professor
prof_agg = df_reviews.groupby(
    ['prof_name','slug','target_rating']
).agg(
    num_reviews = ('review_text','size'),
    num_courses = ('course','nunique'),
    avg_grade_pt = ('grade_pt','mean'),
    text_blob = ('review_text', lambda L: " ".join(L))
).reset_index()
prof_agg.head()
```

	prof_name	slug	target_rating	num_reviews	num_courses	avg_grade_pt	text_blob
0	A.Antony	antony	1.00	1	1	2.000000	By far the worst professor I've ever had, and...
1	A.Kongwatsi	kwagwatsi	3.50	2	2	3.750000	DO NOT TAKE PYS-Coder "Mastermind Social Cogni...
2	A.Sharma	sharma_s	1.80	5	1	3.400000	Hey honey, it's hard to maintain your face...
3	A.U. Shanker	shanker_a.u.	2.44	25	2	3.361111	Lectures are pretty dry and difficult to follo...
4	Aaron Allen	allen_aaron	5.00	1	0	NaN	Great at teaching material, very funny if you...

Next steps: [Generate code with prof\\_agg](#) [View recommended plots](#) [New interactive sheet](#)

```
#Compute NLP features
import re
from textblob import TextBlob
from nltk.sentiment.vader import SentimentIntensityAnalyzer
import nltk
nltk.download('vader_lexicon')

def nlp_feats(text):
    blob = TextBlob(text)
    tokens = re.findall('[\w']+', text.lower())
    return pd.Series({
        'tk_sentiment': blob.sentiment.polarity,
        'vocab_size': len(set(tokens)),
        'avg_word_len': (sum(len(t) for t in tokens)/len(tokens)) if tokens else 0
    })

nlp_df = prof_agg['text_blob'].apply(nlp_feats)

#vADER features
sia = SentimentIntensityAnalyzer()
df_reviews['vader_compound'] = df_reviews['review_text'].apply(lambda txt: sia.polarity_scores(txt)['compound'])
vader_df = df_reviews.groupby('prof_name')['vader_compound'].mean().reset_index().rename(columns={'vader_compound':'vader_sentiment'})

df_prof = pd.concat([prof_agg.drop(columns='text_blob').reset_index(drop=True), nlp_df, axis=1)
df_prof = df_prof.assign(vader_df_0=prof_name, how='left')
df_prof.head()
```

Downloading package vader\_lexicon to /root/nltk\_data...

	prof_name	slug	target_rating	num_reviews	num_courses	avg_grade_pt	tk_sentiment	vocab_size	avg_word_len	vader_sentiment
0	A.Antony	antony	1.00	1	1	2.000000	-0.100000	43.0	5.424783	-0.409300
1	A.Kongwatsi	kwagwatsi	3.50	2	2	3.750000	0.09524	184.0	5.507376	0.191100
2	A.Sharma	sharma_s	1.80	5	1	3.400000	-0.19118	130.0	5.024422	-0.597300
3	A.U. Shanker	shanker_a.u.	2.44	25	2	3.361111	0.050501	629.0	5.127182	0.188004
4	Aaron Allen	allen_aaron	5.00	1	0	NaN	0.185000	22.0	5.830384	0.874800

Next steps: [Generate code with df\\_prof](#) [View recommended plots](#) [New interactive sheet](#)

```
#drop missing target and filter >= 5 reviews
df_prof = df_prof.dropna(subset=['target_rating'])
df_prof = df_prof[df_prof['num_reviews'] >= 5].reset_index(drop=True)

#annotate remaining NaNs
for c in ['avg_grade_pt','tk_sentiment','vocab_size','avg_word_len','vader_sentiment']:
    df_prof[c].fillna(df_prof[c].median(), inplace=True)

print(f"Cleaned professor table shape: ", df_prof.shape)
df_prof.head()
```

Cleaned professor table shape: (1808, 10)

FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy. For example, when doing df[col].method(value, inplace=True), try using df.method(col: value, inplace=True) or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

	prof_name	slug	target_rating	num_reviews	num_courses	avg_grade_pt	tk_sentiment	vocab_size	avg_word_len	vader_sentiment
0	A.Sharma	sharma_s	1.8000	5	1	3.400000	-0.131118	130.0	5.025422	-0.557200
1	A.U. Shanker	shanker_a.u.	2.4400	25	2	3.361111	0.050501	629.0	5.127182	0.188004
2	Aaron Barrett	barrett	2.1667	6	4	3.540000	0.141905	187.0	5.118407	0.268467
3	Aaron Finkel	finkel	3.2000	30	2	3.141379	0.060764	607.0	5.080369	0.411567
4	Aaron Kjelasson	kjelasson_aaron	3.4444	9	2	3.637500	0.120761	293.0	5.017241	0.438700

Next steps: [Generate code with df\\_prof](#) [View recommended plots](#) [New interactive sheet](#)

```
#how many professors do we have?
num_professors = len(all_details)
print(f"Number of professors: {num_professors}")

#how many total reviews?
num_reviews = len(df_reviews)
print(f"Number of reviews: {num_reviews}")

#what's in the DataFrame?
print(f"DataFrame shape: ", df_reviews.shape)
print(f"Columns: ", list(df_reviews.columns))

df_reviews.head()
```

1

Number of professors: 13427  
Number of reviews: 37533

DataFrame shape: (37533, 8)  
Columns: ('prof\_name', 'slug', 'target\_rating', 'course', 'review\_text', 'grade\_pt', 'vader\_compound')

	prof_name	slug	target_rating	course	review_text	grade_pt	vader_compound
0	Abeni Pradhan	pradhan	3.8867	None	Pretty funny guy and he knows his stuff too. P...	B+	3.3
1	Abeni Pradhan	pradhan	3.8867	None	He's a pretty funny guy and he knows his stuff too. P...	B-	2.7
2	Abeni Pradhan	pradhan	3.8867	NFBC12	Didn't learn anything, but don't need to. Eas...	A	4.0

1

2

```
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_absolute_error
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, r2_score
```

2

3

```
#Features and target
feature_cols = ['num_reviews', 'num_courses', 'avg_grade_pt', 'tb_sentiment', 'vocab_size', 'avg_word_len', 'vader_se :
x = df_prof[feature_cols]
y = df_prof['target_rating']

X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=42)
```

3

4

```
models = [
    "KNN": Pipeline([("impute", SimpleImputer(strategy="median")), ("scale", StandardScaler()), ("knn", KNeighborsRegr :
    "RF": Pipeline([("impute", SimpleImputer(strategy="median")), ("scale", StandardScaler()), ("rf", RandomForest :
    "LR": Pipeline([("impute", SimpleImputer(strategy="median")), ("scale", StandardScaler()), ("lr", LinearReg :
    )

results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    results.append(name, mean_absolute_error(y_test, preds), r2_score(y_test, preds))
print(f'{name}: MAE={results[1][1]:.3f}, R2={results[1][2]:.3f}')

KNN: MAE=0.348, R2=0.880
RF: MAE=0.333, R2=0.838
LR: MAE=0.324, R2=0.828
```

4

5

```
import matplotlib.pyplot as plt

preds = model.fit("LR").predict(X_test)
plt.scatter(y_test, preds, alpha=0.3)
plt.plot([1,5],[1,5], 'k--')
plt.xlabel('True Rating')
plt.ylabel('Predicted Rating')
plt.title('True vs Predicted (Linear Regression)')
plt.show()
```

5

6

```
import matplotlib.pyplot as plt
import numpy as np

model = model.fit("RF").named_steps['rf']
feature_names = X_train.columns

importances = model.feature_importances_
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(8, 6))
plt.barh([feature_names[i] for i in indices],
         importances[indices],
         align='center')
plt.title('Feature Importances')
plt.xlabel('Importance')
plt.gca().invert_yaxis()
plt.show()
```

6

7

```
import numpy as np
import matplotlib.pyplot as plt

lr_model = model["LR"].named_steps['lr']
coef = lr_model.coef_
features = X_train.columns

idx = np.argsort(np.abs(coef))[::-1]

plt.figure(figsize=(6,3))
plt.barh([features[idx[i]] for i in range(5)], coef[idx][5])
plt.title('Top 5 Linear Regression Coefficients')
plt.xlabel('Coefficient value')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()
```

7

8

```
from sklearn.inspection import permutation_importance

perm = permutation_importance(
    model["KNN"], X_test, y_test,
    n_repeats=10,
    random_state=42,
    n_jobs=-1
)

importances = perm.importance_mean
features = X_test.columns
idx = np.argsort(importances)[::-1]

plt.figure(figsize=(6,3))
plt.barh([features[idx[i]] for i in range(5)], importances[idx][5])
plt.title('Top 5 KNN Permutation Importances')
plt.xlabel('Decrease in R^2 when permuted')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

8

9

```
import matplotlib.pyplot as plt

decrease_in_r2_when_permuted = perm.importance_mean
features = X_test.columns
idx = np.argsort(importances)[::-1]

plt.figure(figsize=(6,3))
plt.barh([features[idx[i]] for i in range(5)], decrease_in_r2_when_permuted[idx][5])
plt.title('Top 5 KNN Permutation Importances')
plt.xlabel('Decrease in R^2 when permuted')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

9